

Visualization & Dashboard Analysis Report

Online Retail II — Power BI Two-Page Dashboard

Prepared by Sivaan · Data Analytics Portfolio Project

This report documents and critically evaluates the design and analytical value of a two-page Power BI dashboard built on the Online Retail II dataset. All findings, patterns, and figures referenced below are drawn directly from the two dashboard screenshots supplied for this review. Where a statement goes beyond what is literally visible — an inference about chart construction, or a caveat about what cannot be confirmed — it is explicitly flagged as such. The dashboard's screenshots show no slicers or interactive filter controls on either page; this report does not describe or imply any interactivity beyond what is visible.

Contents

1. Dashboard Overview
2. KPI Visualization Analysis
3. Page 1 — Monthly Revenue Trend
4. Page 1 — Orders, Cancellations & Returns
5. Page 1 — Weekly Trend Analysis
6. Page 1 — Overall Analytical Value
7. Page 2 — Monthly Net Revenue vs Returns
8. Page 2 — Monthly Active Customers
9. Page 2 — Geographic Revenue Distribution
10. Visualization Reasoning
11. Dashboard Strengths
12. Dashboard Limitations
13. Business Insights Supported by the Dashboard
14. Recommended Improvements
15. Final Dashboard Assessment

1. Dashboard Overview

The dashboard under review is a two-page Power BI report built on the Online Retail II transaction dataset. Its overall purpose is to give a business stakeholder — without opening the underlying data — a fast, visual read on how sales, returns, and customer activity have moved over time, and where revenue is geographically concentrated. Both pages carry the same header text, “Revenue & Customer Activity Dashboard,” rendered in a dark theme with a consistent colour palette (blue for revenue, gold/orange for gross figures and cancellations, teal for orders and active customers, tan for the dominant country), which gives the two pages a unified visual identity even though their analytical focus differs.

Page 1 — Revenue & Customer Activity Dashboard

Page 1 functions as the executive-level revenue and customer activity overview. Reading left to right, top to bottom, it presents:

- A vertical stack of four KPI cards (Gross Sales, Net Revenue, Returns, Customers) anchoring the page on the left.
- A monthly trend chart comparing Gross Sales and Net Revenue (top-centre).
- A monthly trend chart of Orders, Cancellations, and Returns (top-right).
- A weekly trend chart combining Gross Sales, Net Revenue, Orders, and Returns (bottom).

This combination — a KPI summary plus three complementary trend views at different granularities — positions Page 1 as the page a manager would look at first: it establishes “where do we stand” (the cards) and “how did we get here” (the trends), before any breakdown by geography is introduced.

Page 2 — Revenue, Customer Activity & Geographic Analysis

Page 2 goes one layer deeper on two dimensions Page 1 only summarised — customer activity and geography — while re-examining revenue against returns from a different angle:

- A dual-axis line chart plotting Net Revenue and Returns by month (top), with Active Customer counts embedded in the axis labels.
- A line chart of Active Customers over time (bottom-left).
- A pie chart showing the distribution of Gross Sales by country (bottom-right).

Where Page 1 answers “what happened to revenue and activity,” Page 2 begins to answer “who was buying, and where,” which is a natural next question once the top-line trend has been established.

2. KPI Visualization Analysis

Page 1's four KPI cards sit in a single vertical column, which reads as a deliberate top-to-bottom sequence: Gross Sales, Net Revenue, Returns, Customers.

Gross Sales — 20.97M GBP

Gross Sales represents total sales value before any deduction for returned or cancelled orders — the full value of everything sold, regardless of whether some of it later came back. It answers the most basic business

question, “how much did we sell?”, and functions as the top-line reference point every other revenue figure on the dashboard is measured against.

Net Revenue — 19.45M GBP

Net Revenue is Gross Sales minus the value of returns and cancellations — the revenue actually realised once returned goods are accounted for. It is the more representative figure for understanding true business performance, because Gross Sales alone can overstate how much revenue the business actually retained. Net Revenue answers “how much did we actually keep?” rather than “how much did we ring up?”

Returns — 1.53M GBP

Showing Returns as its own card, rather than folding it silently into Net Revenue, matters because a return figure hidden inside a net number is invisible to a stakeholder scanning the dashboard — they would see 19.45M and have no way to know whether that reflects a business with negligible returns or one absorbing a large and possibly rising cost. Displaying Returns separately answers “how much value came back to us, and is that a small or material share of what we sold?” — turning it into something that can be watched, benchmarked, and investigated on its own.

Customers — 5.94K (unique)

The Customers card measures the count of unique customers transacting in the reporting period. Revenue figures alone say nothing about how many people generated that revenue — the same 19.45M in Net Revenue could come from a small base of high-spending customers or a much larger base of light spenders, and those are very different businesses to manage. Placing a customer-activity metric alongside the revenue KPIs answers “how broad is the customer base behind these numbers?” and sets up the customer-activity charts that follow.

The Gross Sales → Returns → Net Revenue relationship

Read together, the three revenue cards tell a small, self-contained story: Gross Sales (20.97M) is reduced by Returns (1.53M) to arrive at Net Revenue (19.45M) — 20.97M minus 1.53M equals 19.44M, consistent with the displayed 19.45M figure within normal rounding. Presenting all three side by side, rather than Net Revenue alone, gives the viewer the full picture in one glance: not just the destination number, but the size of the adjustment that produced it. A dashboard showing only Net Revenue would communicate the outcome without the context needed to judge whether returns are a minor factor or a meaningful drag on performance; showing all three lets a viewer form that judgement directly from the KPI row, without digging into an underlying report.

3. Page 1 — Monthly Revenue Trend

The “Monthly Trend Analysis For Gross and Net Sales” visual is an area/line combination chart plotting Sum of NetRevenue and Sum of GrossSales across the months January through December, grouped under quarter labels (Qtr 1–Qtr 4).

A line chart, with the area beneath it shaded, is an appropriate choice here because the underlying analytical question is fundamentally about change over time — how sales moved from one month to the next — and a line's connected points are what let the eye trace the shape of that change: rising, falling, flat, or turning at a specific point. The chart lets the viewer make one comparison directly: how closely Gross Sales and Net Revenue track each other across the year, and where the gap between them widens or narrows. Because both series

share the same axis and are plotted together, the vertical distance between the gold Gross Sales line and the blue Net Revenue line at any given month is a visual proxy for that month's returns.

Visible pattern

Across the twelve months shown, both lines sit close together throughout, with Gross Sales consistently at or just above Net Revenue — consistent with the modest overall gap implied by the two KPI cards (20.97M vs. 19.45M). The shape of the trend is the more striking feature: both lines are relatively flat, sitting toward the lower end of the y-axis (roughly the 1.0M–1.5M band), from January through August, then rise steadily from September through November, reaching the chart's highest point in November before dipping slightly in December. This is a classic seasonal-retail shape — a slow first three quarters followed by a fourth-quarter peak — though whether that peak is a genuine calendar-driven seasonality effect or partly an artefact of how the underlying data was aggregated cannot be fully confirmed from this chart alone.

A structural note on this chart

Unlike the Orders/Cancellations/Returns chart and the Weekly Trend chart on the same page — both of which plot a continuous timeline running from Jan 2010 through at least Jul 2011 — this chart's x-axis lists each calendar month exactly once (January through December), with no year attached. That strongly suggests the chart is aggregating by calendar month across the full multi-year date range rather than showing a single continuous month-by-month timeline. This is a reasonable design choice for surfacing a recurring seasonal pattern, but it also means a month such as December is likely combining more than one calendar year's December into a single point, which can mask year-over-year growth or decline. This is an inference about the chart's construction based on its axis structure, not something the screenshot states outright, and is flagged here as a design characteristic worth confirming rather than a confirmed fact.

4. Page 1 — Orders, Cancellations & Returns

This visual — “Monthly Trend Analysis For Orders, Cancellations and Returns” — plots three series (Sum of Cancellations, Sum of Orders, Sum of Returns) across a continuous timeline running from roughly Jan 2010 through Jul 2011, and, based on the density of points past the last labelled tick, likely somewhat beyond it.

What this visual adds beyond the revenue chart

The Monthly Revenue Trend chart (Section 3) shows how much money moved; this chart shows how much activity produced that money and how much of it was later reversed. Revenue and operational activity do not necessarily move together — a business could see stable revenue while order volume and cancellations both rise, for instance, if average order value falls as order count grows. Placing operational counts on their own chart, rather than only showing their dollar impact, lets a viewer sanity-check the revenue story against the underlying transaction volume that produced it.

Visible pattern

Sum of Orders (the blue line) is by far the most volatile series on this chart, swinging repeatedly between roughly 0.05M and points approaching 0.2M, with a particularly pronounced spike visible in the Jan 2011–Jul 2011 window. Sum of Cancellations (gold) and Sum of Returns (teal) sit much closer to the bottom of the chart for most of the period — visually small relative to Orders — but each shows at least one clear departure from

that baseline: a distinct spike in the Returns series is visible later in the timeline, standing well above the surrounding points and reaching a magnitude comparable to some of the Orders peaks.

Questions this visual raises (not answers)

Because Orders, Cancellations, and Returns are plotted together but not shown as a ratio or overlaid against a shared reference point, this chart prompts rather than resolves several diagnostic questions: Are the months with the highest order volume also the months with the highest returns, or do the two move independently? Do cancellations cluster around particular periods, or are they evenly spread? Is the visible late-period spike in Returns associated with the same period as a spike in Orders, or does it stand apart from order volume? The chart makes these questions easy to ask by putting all three series in view at once, but answering them with confidence — and especially attributing a cause to any one spike — would require a level of data access this screenshot alone does not provide. Nothing in this visual establishes a causal link between order volume and return or cancellation volume.

5. Page 1 — Weekly Trend Analysis

The bottom visual on Page 1 plots Sum of GrossSales, Sum of NetRevenue, Sum of Orders, and Sum of Returns at weekly granularity across a continuous “Year” axis (Jan 2010 through at least Jul 2011), with the plotted lines continuing to the right edge of the chart beyond the last labelled tick, suggesting the underlying data extends further than the axis labels shown.

Why weekly aggregation is useful

Monthly aggregation, as used in Sections 3 and 4, smooths out short-term fluctuation by design — a single bad week gets folded into a month that otherwise looks normal. Weekly aggregation trades that smoothness for resolution: it is the right granularity for spotting a specific short-lived event (a promotion, a stock-out, a processing backlog) that a monthly view would average away entirely. This is why the same underlying figures reappear here at a finer time slice rather than being redundant with the monthly views elsewhere on the page.

What becomes visible at this resolution

The most immediately striking feature of this chart is a sharp, isolated spike — most visible in the Orders series — that rises well above the surrounding weekly values, situated in the vicinity of the Jan 2011 tick mark. Outside of that spike, all four series show the kind of tight, saw-tooth week-to-week volatility that is normal and expected at weekly granularity, oscillating within a fairly narrow band for most of the observed period. Toward the right-hand portion of the chart, Gross Sales and Net Revenue appear to trend upward together, broadly consistent with the fourth-quarter rise seen in the Monthly Revenue Trend chart, though the weekly chart's noisier baseline makes that rise harder to isolate with confidence at a glance compared with the smoother monthly view.

Why this matters operationally

A single anomalous week like the one visible here is exactly the kind of signal a monthly-only dashboard would suppress. Whether that spike reflects a genuine one-time event or a data artefact cannot be determined from this chart alone — it can only be flagged as a specific week worth investigating in the underlying data, which is precisely the operational-monitoring role a weekly view is meant to serve.

6. Page 1 — Overall Analytical Value

Taken together, the four visuals on Page 1 let a stakeholder walk away with a working understanding of recent business performance without opening the source dataset.

Revenue

The KPI cards establish the headline numbers (20.97M gross, 19.45M net), and the Monthly Revenue Trend chart shows that this revenue was not earned evenly across the year: it was concentrated in a rise through the final quarter, following a comparatively flat first eight months.

Customer Activity

The Customers KPI (5.94K unique customers) gives a sense of the size of the customer base, and the Orders series on both the monthly and weekly charts shows that order volume is highly variable week to week and month to month rather than steady — activity arrives in bursts rather than as a smooth stream.

Returns & Cancellations

The Returns KPI (1.53M) quantifies the scale of reversed revenue relative to the 20.97M gross figure, and the Orders/Cancellations/Returns chart shows that, while returns and cancellations sit well below order volume for most of the period, they are not negligible everywhere — at least one period shows a return spike large enough to stand out clearly against the rest of the series.

Time

Between the monthly and weekly views, Page 1 supports two different kinds of time-based reading: the monthly charts are suited to spotting the broad seasonal shape of the year, while the weekly chart is suited to spotting a specific short-lived anomaly that the monthly view would average away. Together they give both a strategic and an operational read on the same underlying activity.

7. Page 2 — Monthly Net Revenue vs Returns

The top visual on Page 2, titled “NetRevenue by ActiveCustomers and Month,” is a dual-axis line chart comparing Sum of NetRevenue (left axis, “NetRevenue”) against Sum of Returns (right axis, “Returns Sa...”, truncated in the screenshot). Showing Returns as its own line on a secondary axis, rather than only as a subtraction baked into a single net figure, lets a viewer visually compare the shape of the two series independently: whether Returns rise and fall in step with Net Revenue, or move on their own schedule. As on Page 1, Net Revenue sits noticeably above Returns for the full period shown, and Returns appear comparatively flat and low relative to Net Revenue's movement, with one visible bump above its otherwise low, steady band roughly a third of the way across the chart.

Any period where the two lines move in noticeably different directions is the kind of point that would deserve closer investigation — for instance, a period where Returns rise while Net Revenue is flat or falling would be more concerning than one where both simply move up together — but distinguishing those cases with confidence is difficult on this particular chart, for the reason below.

A significant, visible design issue

The x-axis of this chart combines three pieces of information into a single compound label at each point: a month name, a year, and an Active Customers count (for example, “December 2011 / 686,” “January 2010 / 783,” and so on). Two problems compound here. First, the month-year combinations do not run in chronological order — the sequence of labels jumps between years unpredictably rather than proceeding in date sequence. Second, and more tellingly, the Active Customers numbers printed beneath each month-year label increase in an almost unbroken ascending sequence from left to right across the entire axis. Taken together, this strongly suggests the axis has been sorted by the Active Customers value — a side effect of how the Month/Year hierarchy and a numeric field were combined in the visual's field configuration — rather than by date.

The practical consequence is serious: a line chart's entire value comes from the assumption that adjacent points represent adjacent points in time, so the shape of the line reads as a trend. If the x-axis is not actually in date order, the line connecting these points does not represent a time trend at all — it connects points that happen to sit next to each other because of their Active Customers value, not because they occurred in sequence. This materially undermines the chart's usefulness for trend-reading in its current form and is the single most significant design issue found on either page of this dashboard.

8. Page 2 — Monthly Active Customers

The bottom-left visual, “Sum of ActiveCustomers by Year, Quarter, Month and Day,” is a line chart plotting the count of active customers over a continuous, correctly chronological time axis (Jan 2010 through at least Jul 2011, ticks at roughly six-month intervals), with the y-axis running from 0 to above 1500.

What it represents

Active Customers counts the number of distinct customers who transacted within each period — a measure of the breadth of the customer base actively engaging with the business, as distinct from how much they spent (revenue) or how many transactions they generated (orders). A business can see rising revenue with a shrinking active-customer base if existing customers are simply spending more, or falling revenue with a stable customer base if average spend per customer drops. The three metrics are related but not interchangeable, and this chart is what allows the “how many people” question to be examined on its own.

Visible pattern

The line is highly volatile rather than smoothly trending: it rises and falls repeatedly across the period, with a pronounced peak visible around Jan 2011 (climbing distinctly above the surrounding months), a decline afterward, and a second, even higher climb later in the series, before a sharp drop at the very right edge of the chart. Whether that final drop reflects a genuine falloff in customer activity or simply an incomplete final period in the underlying data cannot be determined from this chart alone.

Why the three metrics should not be assumed to move together

This chart, read alongside the revenue charts on Page 1 and the top chart on this page, is a useful check against over-interpreting revenue movement as customer-driven, or vice versa. A rise in Net Revenue does not, by itself, confirm that more customers were buying — it could equally reflect the same customers buying more per visit — and this Active Customers chart is the only visual on either page that isolates customer count as its own signal, independent of what those customers spent.

9. Page 2 — Geographic Revenue Distribution

The bottom-right visual, “Distribution of Gross Sales By Countries,” is a pie chart with a legend listing eight countries: United Kingdom, EIRE, Netherlands, Germany, France, Australia, Spain, and Switzerland. Four slice labels are visible on the chart itself: 17.87M (86.35%) on the dominant tan slice, and three smaller labels — 0.68M (3.21%), 0.55M (2.68%), and 0.07M (0.34%) — attached to slices too small to identify individually with confidence from the screenshot alone.

Which market dominates, and what that means

The United Kingdom slice fills the overwhelming majority of the pie — visibly, and by its labelled value of 86.35% of Gross Sales — making it unambiguous that this business's revenue is concentrated in a single home market, with every other country shown contributing a small fraction each. This kind of concentration matters strategically: a business this dependent on one market carries a corresponding concentration risk, where a downturn, regulatory change, or competitive shift in that one country would have an outsized effect on total revenue. It also indicates that international reach, if it exists at all, has not yet meaningfully diversified where revenue actually comes from.

A critical evaluation of the pie chart choice

A pie chart does succeed at the one thing it is being asked to do here: communicate, at a glance, that one country dominates. Where it struggles is precisely the comparison a stakeholder would want to make next — how do the remaining seven countries compare to each other? Because pie slices are compared by angle and area, and human perception of angle and area differences is far less precise than perception of position or length, the seven non-UK slices are compressed into a thin wedge of the circle that is difficult to distinguish without relying on the labels and legend colours. Two countries with meaningfully different revenue can end up looking almost identical as pie slices. A sorted horizontal bar chart — countries ranked by Gross Sales, longest bar at top — would preserve the same overall story (UK far ahead of everyone else) while making the relative ranking and magnitude of the remaining seven countries immediately and precisely readable, which the pie chart in its current form does not achieve.

10. Visualization Reasoning

The following applies a reasoning-first framework — analytical question, chart-type rationale, what the visual reveals, how it could be used, and its limitation — to each of the dashboard's major visuals, rather than describing chart appearance alone.

Monthly Revenue Trend (Gross vs Net Sales)

Analytical Question: How did total sales and realised revenue move across the year, and how large is the gap between them?

Visualization Type: Line/area chart — time is continuous and the shape of change, not just individual values, is the point.

What It Reveals: A flat January–August period followed by a Q4 rise peaking in November, with Gross and Net Sales tracking closely together throughout.

Business Use: Flags Q4 as the period driving most of the year's sales — useful for planning inventory, staffing, or marketing timing around that window.

Limitation: Likely aggregates by calendar month across multiple years, which can hide year-over-year change and blend uneven month coverage into a single point.

Orders, Cancellations & Returns (Monthly)

Analytical Question: How does transaction volume and its reversal (cancellations, returns) vary month to month?

Visualization Type: Multi-series line/scatter chart on a continuous timeline — needed to compare three series' shapes against one another.

What It Reveals: Orders are highly volatile; cancellations stay low and steady; returns are mostly low but show one pronounced late-period spike.

Business Use: Signals a specific period where returns broke from their normal pattern, worth flagging for root-cause follow-up.

Limitation: Does not show returns or cancellations as a rate relative to orders, so it cannot by itself say whether the spike is proportionally unusual or simply a byproduct of higher volume that period.

Weekly Trend Analysis

Analytical Question: Are there short-lived spikes or dips that a monthly view would smooth away?

Visualization Type: Multi-series line chart at weekly grain — the finer time unit needed to surface single-week events.

What It Reveals: One sharp, isolated spike (most visible in Orders) standing well above the surrounding weekly volatility.

Business Use: Provides a specific week to investigate operationally, rather than a vague “that quarter was unusual.”

Limitation: High week-to-week noise makes the underlying trend harder to read at a glance than the monthly view; four series in one chart adds visual density.

NetRevenue by ActiveCustomers and Month (Page 2, top)

Analytical Question: How do Net Revenue and Returns move together over time, and how does that relate to customer activity?

Visualization Type: Dual-axis line chart — intended to let two differently-scaled series (Net Revenue, Returns) be compared on the same time axis.

What It Reveals: Net Revenue consistently exceeds Returns, with Returns comparatively flat aside from one visible bump.

Business Use: In principle, a quick visual check on whether returns are rising in step with revenue or moving independently.

Limitation: The x-axis is not in chronological order (it appears sorted by the Active Customers value instead), which undermines the chart's ability to be read as a time trend at all in its current form.

Active Customers Trend (Page 2, bottom-left)

Analytical Question: How has the size of the actively transacting customer base changed over time?

Visualization Type: Single-series line chart on a correctly chronological continuous axis — the standard choice for a single trend over time.

What It Reveals: A volatile pattern with a peak around Jan 2011, a pullback, and a second, higher peak later, before a sharp drop at the final point shown.

Business Use: Lets a stakeholder check whether a revenue change coincided with more customers buying, or the same customers spending differently.

Limitation: A count alone does not show whether the active customers in each period are new or returning, so it cannot distinguish organic growth from repeat-purchase behaviour.

Distribution of Gross Sales by Countries (Pie Chart)

Analytical Question: Which countries generate the business's Gross Sales, and how concentrated is that revenue geographically?

Visualization Type: Pie chart — a reasonable choice for making one thing obvious at a glance: whether a single category dominates the whole.

What It Reveals: The United Kingdom accounts for the overwhelming majority of Gross Sales (86.35%), with every other listed country contributing a small share each.

Business Use: Signals a market-concentration risk and shows that revenue diversification beyond the UK is currently very limited.

Limitation: The seven non-UK slices are too visually compressed to compare precisely against each other — a bar chart would preserve the UK-dominance story while making the smaller countries' relative ranking legible.

KPI Card Row (Gross Sales, Net Revenue, Returns, Customers)

Analytical Question: What are the headline totals for the reporting period, before any time-based or segment-based breakdown?

Visualization Type: Card visuals — the standard choice when a single precise number, not a shape or comparison, is the point.

What It Reveals: The absolute scale of sales, the size of the return-driven adjustment to reach realised revenue, and the size of the customer base behind those figures.

Business Use: Gives an at-a-glance status check without needing to read a single chart — useful as the first thing a time-pressed stakeholder looks at.

Limitation: Cards show no comparison point (no prior-period figure, target, or trend indicator), so a viewer cannot tell from the cards alone whether 20.97M is a good or a concerning number relative to any baseline.

11. Dashboard Strengths

The strongest aspects of the dashboard, and why each is analytically useful rather than merely present:

Clear KPI-first structure: leading with four card visuals gives a viewer the headline numbers before asking them to interpret any chart, a sound design sequence for an executive audience.

Separation of Gross Sales, Returns, and Net Revenue: showing all three rather than only a net figure lets a viewer judge the scale of returns for themselves instead of trusting a single blended number.

Multiple time granularities: pairing monthly and weekly views of the same underlying activity lets the dashboard serve two different reading needs — a strategic seasonal read and an operational anomaly-spotting read — without one crowding out the other.

Inclusion of customer activity alongside revenue: the Customers KPI and the Active Customers trend chart correctly treat “how many customers” as a distinct question from “how much revenue,” preventing the dashboard from being read purely as a revenue report.

Explicit treatment of cancellations and returns: dedicating chart space to Orders, Cancellations, and Returns together, rather than only netting them into a single revenue figure, keeps a real operational cost visible rather than hidden inside an aggregate.

Geographic view: including a country-level breakdown, even in a chart form with room for improvement, adds a dimension — where revenue comes from — that neither page would otherwise cover at all.

Consistent use of trend visualizations: line and area charts are used for every genuinely time-based question in the dashboard, the analytically correct default rather than reaching for chart types not suited to showing change over time.

12. Dashboard Limitations

Presented here as gaps to be aware of, and, where relevant, the analytical layer that is currently missing as a result — not as failures of the work done so far.

No slicers or interactive filtering: both pages are static views — there is no way, from the dashboard itself, to isolate a single month, country, or customer segment for closer inspection. The layer this removes is drill-down: a viewer who notices something worth investigating has no way to filter down to it within the dashboard.

Repeated revenue views across pages: Net Revenue and Returns appear as their own comparison chart on both Page 1 and Page 2 without either page building meaningfully on the other's version of that comparison, using chart space that could otherwise cover a dimension neither page currently addresses, such as product category.

Dense weekly trend visualization: four series plotted together at weekly grain produces a visually busy chart; the short-lived spike it successfully reveals is easy to spot, but routine week-to-week movement underneath it is harder to read precisely than the smoother monthly views.

Complicated date hierarchy on Page 2: as detailed in Section 7, the top chart's x-axis is not in chronological order, which removes this visual's ability to be trusted as a genuine time trend until the sort order is corrected.

Pie chart limitations for country comparison: as detailed in Section 9, the geographic chart makes the UK's dominance clear but does not let a viewer compare the remaining seven countries against each other with any precision.

Limited diagnostic segmentation: every chart on both pages segments the data by time, and on Page 2 by country — none breaks revenue, orders, or returns down by customer segment, price tier, or any other business-relevant grouping. The dashboard can say that something changed in a given period, but not for whom or in which part of the business.

No product/category analysis: neither page includes any visual built around product, SKU, or category — a significant gap for a retail dataset, since a viewer can see that sales rose in Q4 but has no way from this dashboard to know which products drove that rise.

No direct retention or cohort visualization: the Active Customers chart shows how many customers were active in each period, but not whether they are the same customers returning or a continually refreshed set of new ones — the retention/cohort layer is absent.

No direct customer-value segmentation: there is no visual distinguishing high-value from low-value customers, so the dashboard cannot currently support a question like “how much of our revenue depends on a small group of high-spending customers.”

13. Business Insights Supported by the Dashboard

Each item below separates a finding — what the visual directly shows — from an insight — why that finding matters in a business context — and is kept to what the visuals can actually support. No causal explanation is claimed beyond what the dashboard provides evidence for.

Revenue Performance

Finding: Gross Sales for the period totalled 20.97M GBP against Net Revenue of 19.45M GBP, and the Monthly Revenue Trend chart shows both figures rising sharply from a relatively flat first eight months into a Q4 peak around November.

Insight: The business's annual revenue is not earned evenly — a large share of it is concentrated in the final quarter — which matters for how far in advance inventory, staffing, and cash flow need to be planned around that window, and makes the fourth quarter the single highest-leverage period to protect from operational disruption.

Customer Activity

Finding: 5.94K unique customers transacted in the period, and the Active Customers trend chart shows this count fluctuating substantially rather than holding steady, with at least two distinct peaks.

Insight: Because customer count moves independently of the revenue figures shown elsewhere on the dashboard, a period of rising revenue cannot automatically be assumed to reflect more customers buying — it could equally reflect existing customers spending more — and this dashboard alone does not resolve which explanation applies to any specific period.

Returns & Cancellations

Finding: Returns (1.53M GBP) represent a modest but non-trivial share of Gross Sales, and the Orders/Cancellations/Returns chart shows this figure is not evenly distributed — it includes at least one period where returns spike well above their typical level.

Insight: A return figure this visible in the KPI row, combined with a visible spike in the trend chart, is exactly the kind of pattern that justifies asking whether a specific period, product line, or customer segment is disproportionately responsible for returned value — a question this dashboard raises but, without product- or segment-level detail, cannot answer on its own.

Geography

Finding: The United Kingdom accounts for 86.35% of Gross Sales, with the next-largest visible slices in the low single digits.

Insight: This level of concentration in one market means the business's overall performance is, in practice, very close to being the UK market's performance — a downturn or disruption specific to the UK would have an outsized effect on the totals shown elsewhere on the dashboard, and any ambition toward international diversification is not yet reflected in where revenue actually comes from.

14. Recommended Improvements

Priority 1 — High Impact

- Replace the country pie chart with a sorted horizontal bar chart, ranked by Gross Sales, to preserve the UK-dominance story while making the remaining countries' relative ranking precisely readable.
- Fix the date hierarchy on the Page 2 “NetRevenue by ActiveCustomers and Month” chart so the x-axis sorts chronologically rather than, as it currently appears, by the Active Customers value — a prerequisite for the chart to be trustworthy as a trend at all.
- Simplify the weekly trend chart, either by reducing the number of series shown per view or by adding a moving-average line, to make the underlying trend easier to distinguish from week-to-week noise.

Priority 2 — Analytical Depth

- Add a retention/cohort view — for example, active customers broken into new versus returning each period — to answer whether the business is growing its customer base or re-engaging the same customers.
- Add an RFM-style (recency, frequency, monetary) customer segmentation view to identify which customers are driving the bulk of revenue.
- Add a product/category breakdown of sales and returns, since neither page currently includes any product-level view — likely the single biggest gap for a retail dataset of this kind.
- Add a return-rate view (returns as a percentage of gross sales, by month or category) rather than only showing return value in isolation, making it possible to tell whether a return spike is proportionally unusual or simply tracks higher order volume.

Priority 3 — Interactivity

Slicers were not implemented in the current dashboard. Adding them in a future version could let a stakeholder actively explore the questions this static version can only raise, without requiring a new dashboard build for each one:

- A date-range slicer, to isolate a specific month or quarter — such as the Q4 peak or the anomalous week identified in Section 5 — for closer inspection.
- A country slicer, to see the revenue and returns trend for a single market rather than only the aggregated total.
- A product/category slicer (once product-level visuals are added), to isolate which products are driving a given trend.
- A customer-segment slicer (once segmentation is added), to compare high-value and low-value customer behaviour directly.

15. Final Dashboard Assessment

This dashboard succeeds at its core job: giving a stakeholder a fast, visually coherent read on overall sales performance, the scale of returns relative to gross sales, and, on Page 2, a first look at customer activity and geographic concentration, all without needing to open the underlying dataset. The KPI-first structure is sound, the decision to keep Gross Sales, Returns, and Net Revenue visible as three separate figures rather than one blended number is a genuinely good analytical choice, and the pairing of monthly and weekly views gives the dashboard both a strategic and an operational reading of the same activity.

Its weaknesses are concentrated in two places: chart-level design errors — the non-chronological x-axis on Page 2's top chart being the most serious, and the pie chart's compressed comparability being the clearest secondary example — and the missing analytical layers that neither page currently attempts, most notably product/category analysis, customer segmentation, and any form of retention or cohort view. Consistency is otherwise reasonable, as the colour palette and general layout language carry across both pages, though both pages share an identical header title despite covering different analytical ground — a small missed opportunity to signal each page's distinct focus.

Overall, this dashboard currently accomplishes a solid descriptive layer of analysis: it tells a viewer clearly what happened to revenue, returns, and customer activity over the observed period, and where that revenue is concentrated geographically. What it does not yet attempt is the diagnostic layer — explaining why the Q4 peak occurred, why returns spiked in the period identified, or which customers or products are driving the patterns shown. Building that diagnostic layer, through product-level detail, customer segmentation, and corrected, interactive chart mechanics, is the clear next step for this project.